

# String Matching

## 1 Problem Definition

- Given a (long) string  $T$  (over some (finite) alphabet  $\Sigma$ ), and
- A (short) pattern  $P$  over the same alphabet (or subset)
- Find one or all occurrences of  $P$  in  $T$ , if any

For example:  $T$ :acda**bd**dea**bd**de  
 $P$ :bdde

## 2 Terminology

- Input is  $T[1..n]$  and  $P[1..m]$  with  $m \leq n$
- Alphabet  $\Sigma$  can be pretty much anything
- We say a pattern  $P$  occurs with shift  $s$  in  $T$  if

$$0 \leq s \leq n - m \text{ and } T[s + 1, \dots, s + m] = P[1 \dots m]$$

e.g.

$T$ :	$a$	$c$	$d$	$a$	$b$	$d$	$d$	$e$	$a$	$a$	$b$	$d$	$d$	$e$
$P$ :					→	$b$	$d$	$d$	$e$					

Where we've got a shift of 4

$$T[5..8] = P[1..4] - \text{shift } 4$$

$$T[11..14] = P[1..4] - \text{shift } 10$$

## 3 Shifts

We say  $s$  is

- A valid shift if  $P$  occurs with a shift  $s$  in  $T$ , and
- An invalid shift otherwise

In the example, 4 and 10 are valid shifts, all others are not.

**Goal:** Find all occurrences of  $P$  in  $T$  (all valid shifts), as quickly as you can

## 4 More Terminology

Some notation:

- $\Sigma$  is the alphabet, and  $\Sigma^*$  ("Sigma Star") is the transitive closure, i.e., the set of all (finite) words that you can form using characters from  $\Sigma$ , e.g.

$$\Sigma = \{0, 1\} \Rightarrow \Sigma^* = \{\epsilon, 0, 1, 00, 01, 10, 11, 000, \dots\}$$

where  $\epsilon$  denotes the empty word (the word of length 0)

- Length of a string  $x$  is denoted by  $|x|$
- Concatenation of string  $x$  and  $y$  is denoted  $xy$ , and  $xy$ 's length is  $|xy| = |x| + |y|$
- Test  $x \stackrel{?}{=} y$  can be done by comparing character by character starting at the leftmost position, and takes time  $\Theta(t + 1)$  with  $t$  being length of longest string  $z$  with  $z \sqsubset x$  and  $z \sqsubset y$ ; notice the "+1" to handle case  $t = 0$  i.e. no common prefix
- Note that the notation  $z \sqsubset x$  means  $z$  is the prefix of  $x$
- Here  $t$  is the length of  $z$  (the longest match between the two strings)

## 5 The naive algorithm

Input string  $T[1..n]$ , pattern  $P[1..m]$

Idea simple: just slide the pattern along the string and compare at each of the  $n - m + 1$  many positions

Listing 1: Naive Matcher

```

1 for s=0 to n-m do
2   if P[1...m]=T[s+1...s+m] then
3     print "pattern occurs with shift s"
4   end if
5 end for

```

Notice how we write the comparison as if it were a primitive operation(takes constant time), when in reality it isn't

## 6 Example

$T = abaaaddaabaabae, P = aa$

$T: a b a a a d d a a b a a a e$

```

      a a
        a a
          2 a a
            3 a a
              a a
                a a
                  a a
                    7 a a
                      a a
                        a a
                          10 a a
                            11 a a
                              a a

```

Observe: we include overlaps!

Clearly running time  $O((n - m + 1) \cdot m)$

This is tight in the sense that we have worst case instances that need precisely that (consider  $T = a^n$  and  $P = a^m$ ) = we need to compare the full pattern with whatever we find at each of the  $n - m + 1$  possible shifts

Notice that this is

- $\Theta(n^2)$  if  $m \approx n/2$
- $O(n)$  if  $m = O(1)$  or  $m = n - O(1)$
- $n\sqrt{n} = n^{3/2}$  if  $m \approx \sqrt{n}$

This is not particularly good

The biggest problems seems to be that we discard any information about previous comparisons when we move window to the right

## 6.1 Example

An example: consider  $T = \text{aaaaabaaaaa}$  and  $P = \text{aaaa}$

```

0 aaaaabaaaaa
1 aaaaaabaaaaa
2 aaaaabaaaaa
3 aaaaabaaaaa
4 aaaaabaaaaa
5 aaaaabaaaaa
6 aaaaabaaaa
7 aaaaabaaaa

```

First mismatch in line 2, but do we **really** need to do 3,4, and 5?

## 7 The Rabin-Karp algorithm

We introduce a second string-matching algorithm:

- Still not best-possible
- Still same worst-case running time
- Still no "intelligent skipping"
- But in practice typically much better

### 7.1 A useful idea

- Based on representation of substrings of length  $k$  as  $k$ -digit numbers
- That is, if our alphabet has size ten, represent string 61524 as decimal number 61,524
- This alone does of course not help - still need to make as many comparisons
- Extra idea: do calculations modulo some cleverly chosen number
- One of the very few examples where we need to think about "real world" performance
- Would normally assume that we can do arbitrary precision arithmetic in one step
- We will assume  $\Sigma = \{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$
- Any other alphabet size is also possible, for example binary and hex
- Rabin-Karp does some preprocessing, and then the actual matching
- All the "better" string matching algorithms do some sort of preprocessing
- R-K needs  $\Theta(m)$  time for preprocessing, and then worst case  $\Theta((n - m + 1)m)$  for matching - same as naive matcher
- However, much better "on average" - under certain assumptions

## 8 Decimal Representations

- Given a pattern  $P[1..m]$ , denote by  $p$  its decimal value
- Likewise, given string  $T[1..n]$ , let  $t_s$  denote decimal value of length- $m$  substring  $T = [s + 1..s + m]$ , for  $s = 0, 1, \dots, n - m$
- Clearly,  $T[s + 1..s + m] = P$  iff  $t_s = p$
- Therefore,  $s$  is a valid shift  $\Leftrightarrow t_s = p$

What if we:

- Could compute  $p$  in time  $O(m)$ , and
- Could compute all the  $t_s$  values in total time  $O(n - m + 1)$  and
- Could compare  $p$  with any given  $t_s$  in constant time, then
- We'd have a  $O(m) + O(n - m + 1) = O(n)$  algorithm

We can't... or can we?

We can certainly compute  $p$  from  $P[1..m]$  in time  $O(m)$

$$\begin{aligned}
 p &= P[m] + 10P[m-1] + 100P[m-2] + 1000P[m-3] + \dots \\
 &\quad \dots + 10^{m-2}P[2] + 10^{m-1}P[1] \\
 &= P[m] + 10(P[m-1] + 10P[m-2] + 100P[m-3] + \dots \\
 &\quad \dots + 10^{m-3}P[2] + 10^{m-2}P[1]) \\
 &= P[m] + 10(P[m-1] + 10(P[m-2] + 10P[m-3] + \dots \\
 &\quad \vdots \\
 &\quad = P[m] + 10(P[m-1] + 10(P[m-2] + \dots + 10(P[2] + 10P[1]) \dots))
 \end{aligned}$$

↑ that all looks really scary, but it's just str to int conversion if you look at it

That is, use Horner's rule on  $p(x) = \sum_{i=0}^{m-1} P[m-i]x^i$  in  $x = 10$  (compute from the right)

Can do the same for  $t_0$  (from  $T[1..m]$ ), the first  $m$  characters of  $T$

- We still need  $t + 1, t_2, \dots, t_{n-m}$
- We don't want to spend more than  $\Theta(n - m)$  time for that, i.e. constant time for each  $t_i, i \geq 1$

Simple, we can compute "new"  $t_{s+1}$  based on  $t_s$ :

- Consider e.g.

1562745242133  
 1562745242133

i.e., focus on transition for  $t_s = 6274$  to  $t_{s+1} = 2745$  (with  $m=4$ )

- Notice that (when looking at strings) what we're really doing is
  - Pushing the left-most 6 out, and
  - Adding the 5 as new right most character
- This corresponds to the arithmetic operation  $10(6274 - 6 \cdot 1000) + 5$
- More generally,  $t_{s+1} = 10(t_s - 10^{m-1}T[s+1]) + T[s+m+1]$
- Can precompute  $10^{m-1}$  (it's always going to be  $10^{m-1}$ ), then the execution takes constant time

### Summarising

- Thus, we can compute  $p$  and  $t_0$  in time  $\Theta(m)$ , and can compute all  $t_1, \dots, t_{n-m}$  in total time  $\Theta(n - m)$
- We can find all occurrences of  $P$  in  $T$  with  $\Theta(m)$  for preprocessing of  $P$ , and  $\Theta(n - m)$  for the matching

**But wait**

- Surely, if  $m$  is very big, then  $p$  and  $t_i$  are very long numbers
- We wouldn't be able to store them in a computer word each
- Even if we used arbitrarily long integers (through, say, some class or other data structure in some library), it's be unrealistic to assume that we can do the actual arithmetic in constant time

**So what are we going to do**

Answer is quite simple, really: do everything, i.e., all arithmetic and such, modulo some (possibly much smaller) number  $q$ , so that every number fits into one computer word

- Let's return to our example and chose  $q = 123$ ; note  $h = 10^3 \bmod 123 = 16$
- Consider the string 52424537454...
- Suppose we have computed  $t_0$  (corresponds to 5242) as 76 ( $5242 \bmod 123 = 76$ )
- Suppose we want to compute  $t_1$ ; before we just did this:

$$\begin{aligned} t_1 = t_{s+1} &= 10(t_s - 10^{m-1}T[s+1]) + T[s+m+1] \\ &= 10(t_0 - 10^{m-1}T[1]) + T[0+4+1] \\ &= 10(5242 - 1000 \cdot 5) + 4 \\ &= 10 \cdot 242 + 4 = 2420 + 4 \\ &= 2424 \end{aligned}$$

- Observe that  $2424 \bmod 123 = 87$
- Now working modulo  $q$  (recall  $h = 16, t_0 = 76$ )

$$\begin{aligned} t_1 &= (10(t_0 - hT[s+1]) + T[s+m+1]) \bmod q \\ &= (10(76 - 16 \cdot 5) + 4) \bmod 123 \\ &= (10(76 - 80) + 4) \bmod 123 \\ &= (10(-4) + 4) \bmod 123 \\ &= -36 \bmod 123 = 87 \end{aligned}$$

**That seems to work, or does it**

- Consider  $T=62321462338294$ ,  $P=3214$ , and  $q=123$
- $3214 \bmod 123 = 16$
- If we compare the mod-123 values that we get for  $P$  and the substring of  $T$  at shift 2 (3214), then we get a match (which is OK)
- If we compare the mod-123 values that we get for  $P$  and the substring of  $T$  at shift 5 (4623  $\bmod 123 = 72$ ), then we get a mismatch (which is OK)
- If we compare the mod-123 values that we get for  $P$  and the substring of  $T$  at shift 9 (3829), then we get a match (which is not OK)
- In general,  $t_s \neq p \bmod q$  then we know for sure that  $t_s \neq p$  (which is OK)
- However  $t_s \equiv p \bmod q$  does not necessarily imply  $t_s = p$  (which is not OK)

**Ok, this doesn't seem to work that easily**

- It looks as though we can use the text  $t_s \stackrel{?}{\equiv} p \bmod q$  really only as "heuristic"
- If the answer is no then we're fine
- If the answer is yes then we'll have to check in the old fashioned way and may get a false positive
- Hope that it doesn't happen too often, so that the extra cost isn't too high
- Make  $q$  large, so that we don't get "Too many collisions"