

Transport Layer II

1 Reliable data transfer

We will:

- Incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- Consider only unidirectional data transfer - but control info will flow on both directions
- Use finite state machines (FSM) to specify sender, receiver



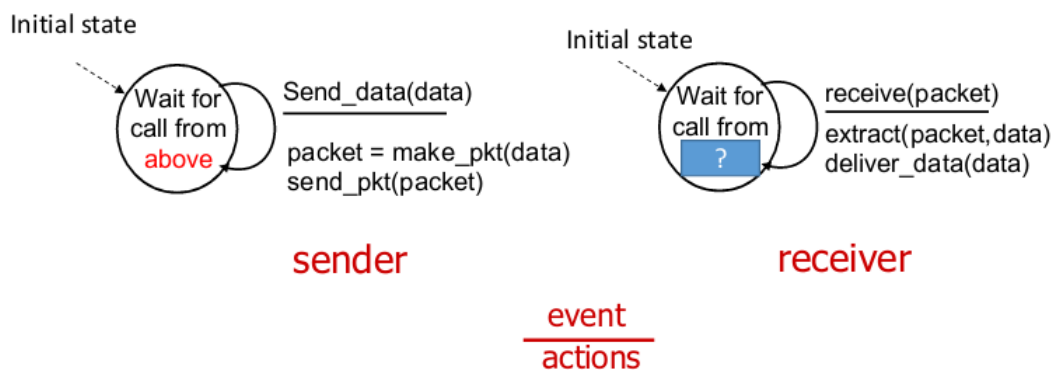
1.1 rdt1.0: reliable transfer over a reliable channel

Underlying channel perfectly reliable:

- No bit errors
- No loss of packets

Separate FSMs for sender, receiver:

- Sender sends data into underlying channel
- Receiver reads data from underlying channel



Important: What layer is this in?

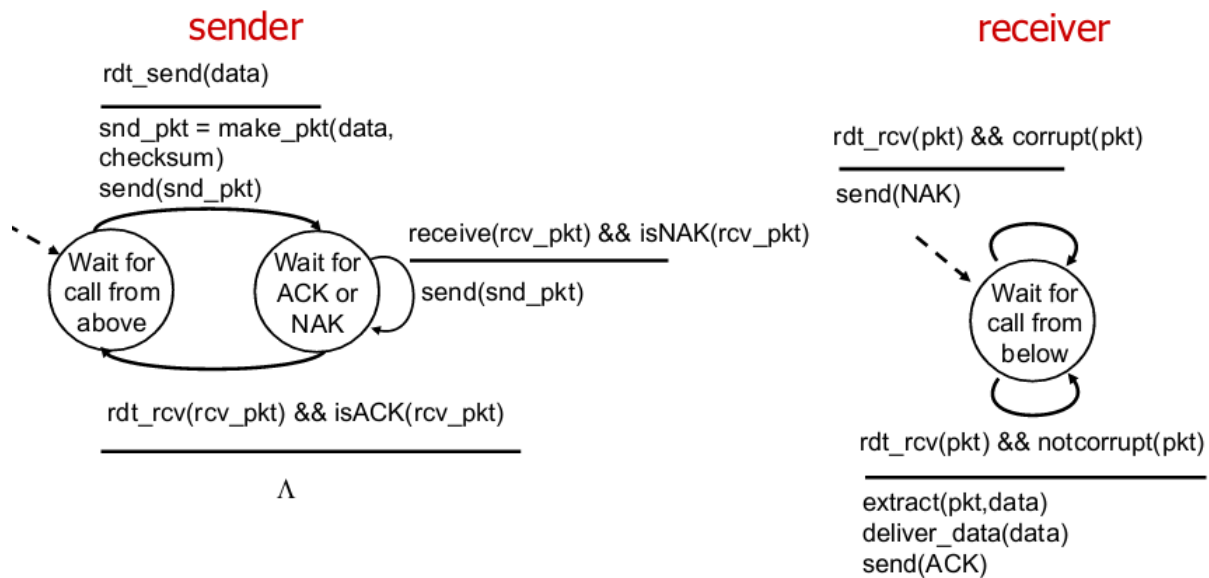
This is happening in the transport layer

1.2 rdt2.0

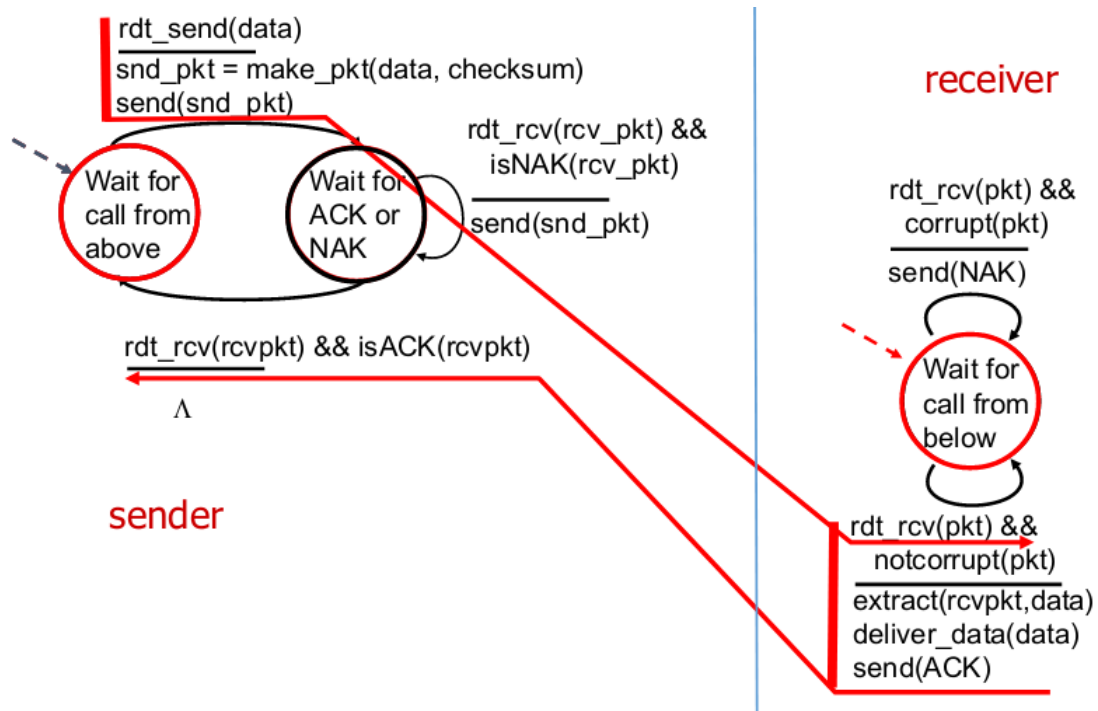
1.2.1 Channel with bit errors

- Underlying channel may flip bits in packet, checksum to detect bit errors
- The question: how to recover from errors:
 - Acknowledgements (ACKs): receiver explicitly tells sender that packet received OK
 - Negative acknowledgements (NAKs): receiver explicitly tells sender that packet had errors
 - Sender retransmits packet on receipt of NAK
 - Using ACKs and NAKs is known as ARQ (Automatic Repeat reQuest) protocols
 - * Error detection. Sender embeds extra bits in packets
 - * Feedback. Receiver provide sender with feedback
 - * Retransmission. Retransmit erroneous packets
- New mechanisms in rdt2.0 (beyond rdt1.0)
 - Error detection
 - Feedback: control msgs (ACK (1), NAK(0)) from receiver to sender

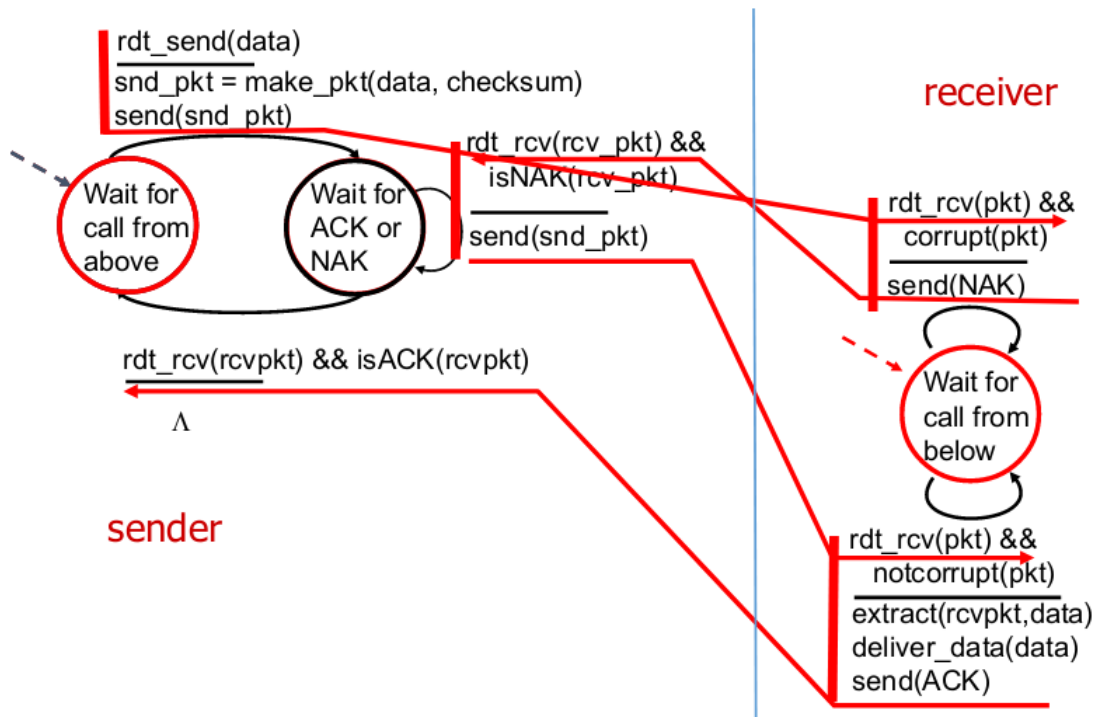
1.2.2 FSM specification



1.2.3 Operation with no errors



1.2.4 Error scenario



1.2.5 Fatal flaw

What happens if ACK/NAK corrupted:

- Sender doesn't know what happened at receiver
- Can't just retransmit: possible duplicate

Handling duplicates:

- Sender retransmits current packet if ACK/NAK corrupted
- Sender adds **sequence number** to each packet
- Receiver discards (doesn't deliver up) duplicate packet

Stop and wait:

- Sender sends one packet, then waits for the receiver response

1.3 rdt3.0

1.3.1 Channels with errors and loss

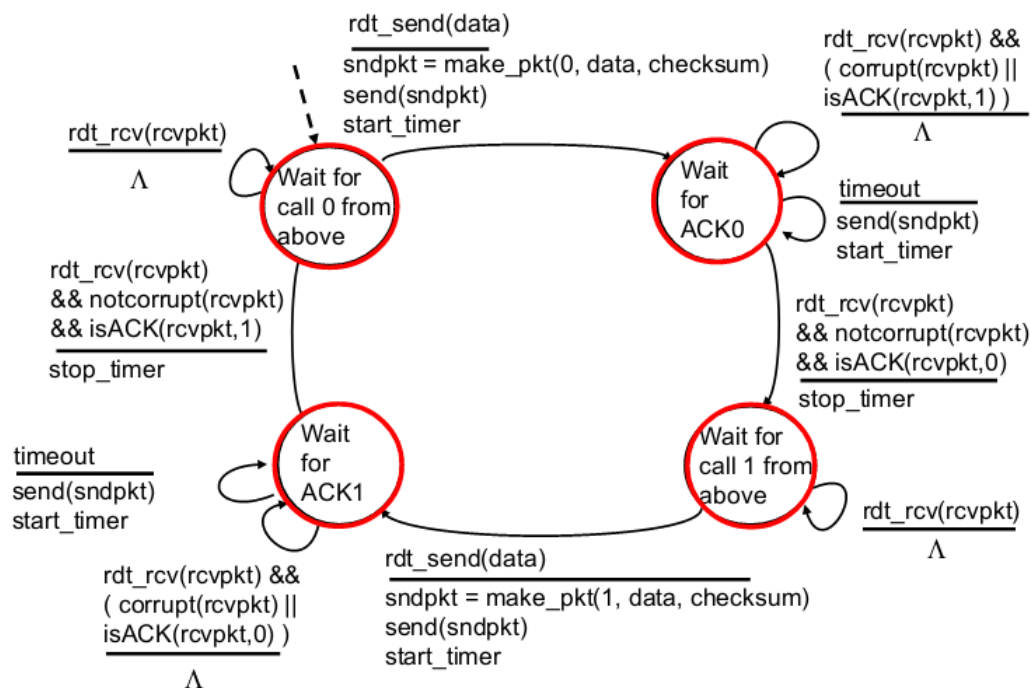
New assumption:

- Underlying channel can also lose packets (data, ACKs)
 - Checksum, seq. #, ACKs, retransmissions will be of help, but not enough

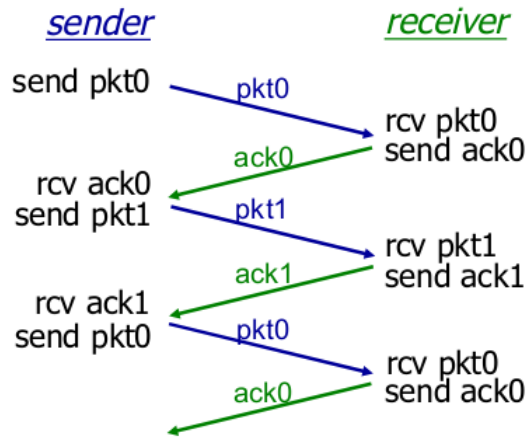
Approach:

- Sender waits "reasonable" amount of time for ACK
 - Retransmits if no ACK received in this time
 - If no packet (or ACK) just delayed (not lost):
 - * Retransmission will be duplicate, but seq. #'s already handles this
 - * Receiver must specify seq # of packet being ACKed
 - Requires countdown timer


1.3.2 Sender



1.3.3 In action




(a) no loss




(b) packet loss

Note: alternating-bit protocol: packet sequence numbers alternate between 0 and 1.

14

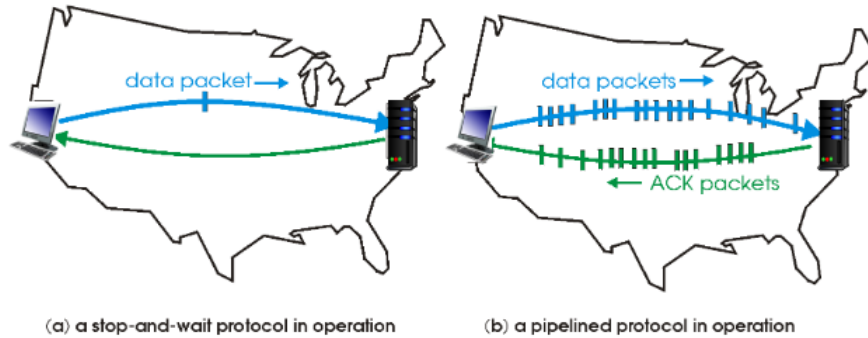


(c) ACK loss



(d) premature timeout/ delayed ACK

2 Pipelined protocols



Pipelining has the following consequences for reliable data transfer protocols:

- Range of sequence numbers must be increased
 - Unique sequence number and there may be multiple, in-transit, unacknowledged packets
- Multiple packet buffering at sender and/or receiver
 - Sender buffers packets that have been transmitted but not yet acknowledged
 - Buffering of correctly received packets
- Range of sequence numbers needed and the buffering requirements will depend on the manner in which a data transfer protocol responds to lost, corrupted, and overly delayed packets

Two generic forms of pipelined protocols:

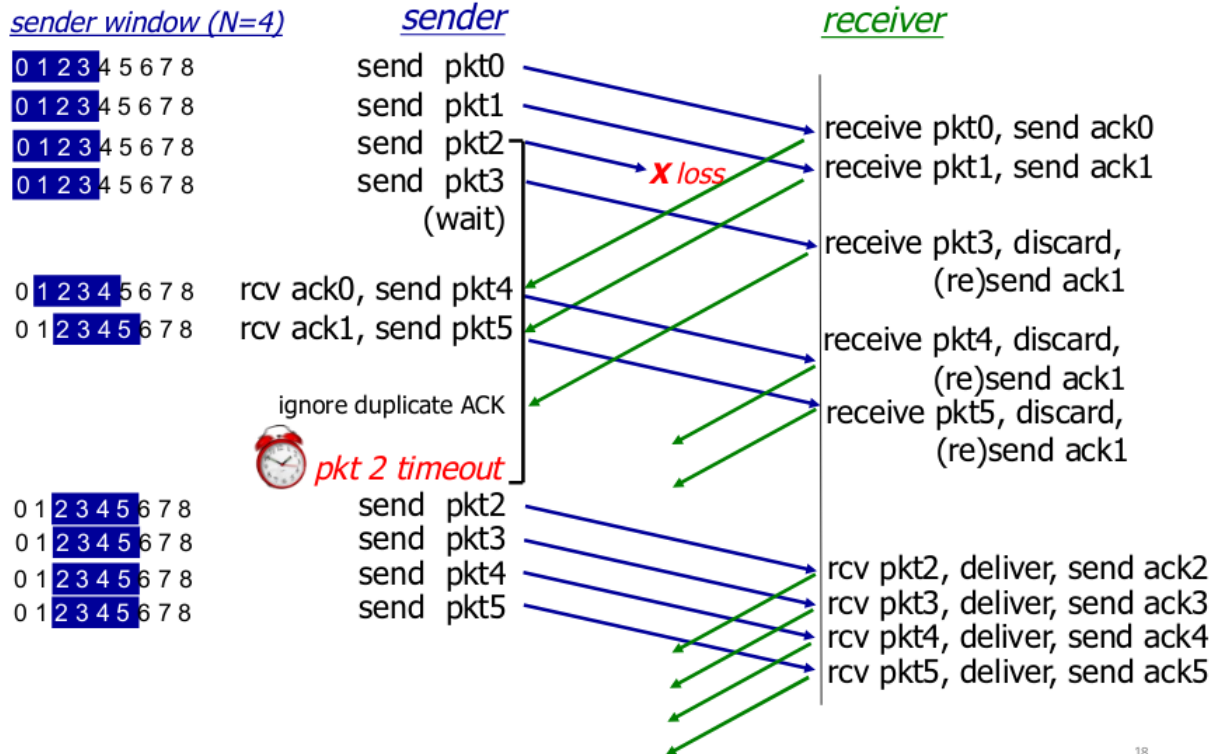
Definition: Go-back-N

- Sender can send multiple packets without waiting for ACK
- Sender can have up to N unacked packets in pipeline
- Receiver only sends cumulative ack, doesn't ack packet if there's a gap (so if it acks packet 4, it means it has packet 1,2,3 and 4)
- Sender has timer for oldest unacked packet, when timer expired, retransmit all unacked packets

Definition: Selective repeat

- Sender can have up to N unacked packets in pipeline
- Receiver sends individual ack for each packet
- Sender maintains timer for each unacked packet, when timer expires, retransmit only that unacked packet

2.1 GBN(Go Back N) in action



2.2 Selective repeat

- Receiver individually acknowledges all correctly receives packets. Buffers packets, as needed, for eventual in-order delivery to upper layer
- Sender only resends packets for which ACK not received. Sender timer for each unACKed packet
- Sender window
 - N consecutive seq #'s
 - Limits seq #'s of sent, unACKed packets

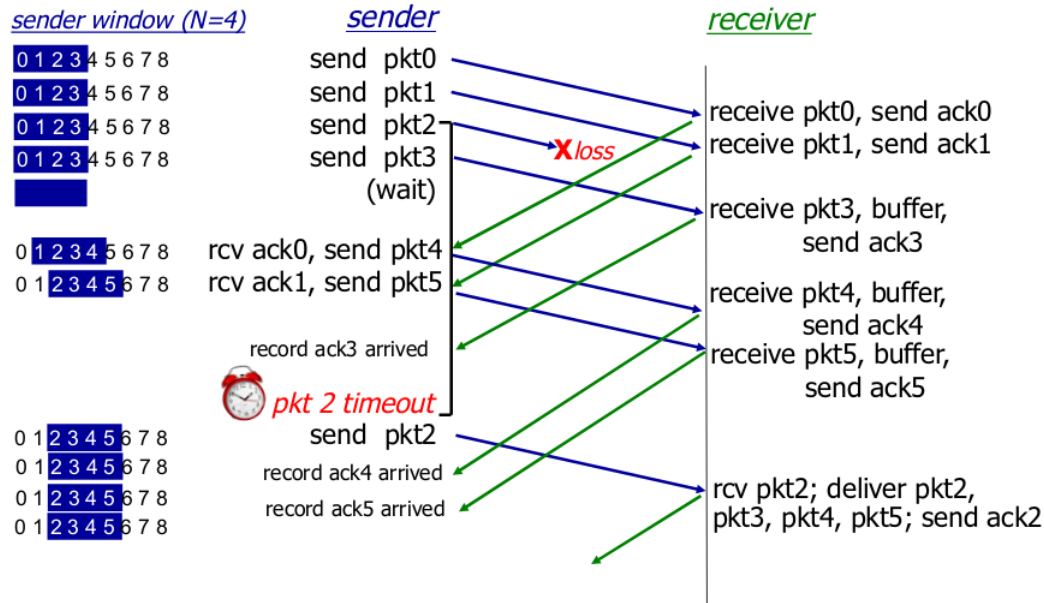
Receiver:

- packet n in $[rcvbase, rcvbase+N-1]$
 - Send ACK(n)
 - Out-of-order: buffer
 - In-order: deliver(also deliver buffered, in-order packets), advance window to next not-yet-received packet
- packet n in $[rcvbase-N, rcvbase-1]$
 - Ack(n)
- Otherwise:
 - Ignore

Sender:

- Data from above:
 - If next available seq # in window, send packet
- timeout(n)
 - Resend packet n , restart timer
- Mark packet n as received
- If n smallest unACKed packet, advance window base to next unACKed seq #


2.2.1 Selective repeat in action



2.2.2 Selective repeat dilemma

Example:

- Finite range of seq # s: 0,1,2,3
- Window size=3
 - Receiver sees no difference in two scenarios
 - Duplicate data accepted as new in (b)



Note: That curtain is there to show that there is a lack of knowledge between the sender and receiver