

Transport Layer II

1 Reliable data transfer

We will:

- Incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- Consider only unidirectional data transfer - but control info will flow on both directions
- Use finite state machines (FSM) to specify sender, receiver



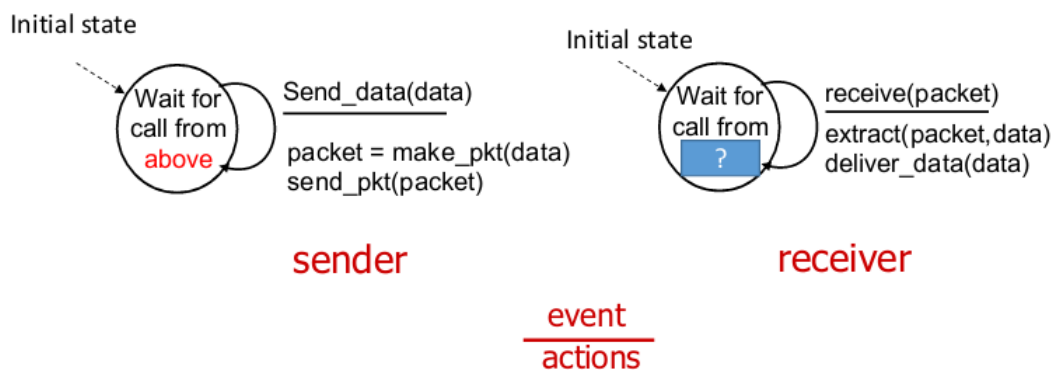
1.1 rdt1.0: reliable transfer over a reliable channel

Underlying channel perfectly reliable:

- No bit errors
- No loss of packets

Separate FSMs for sender, receiver:

- Sender sends data into underlying channel
- Receiver reads data from underlying channel



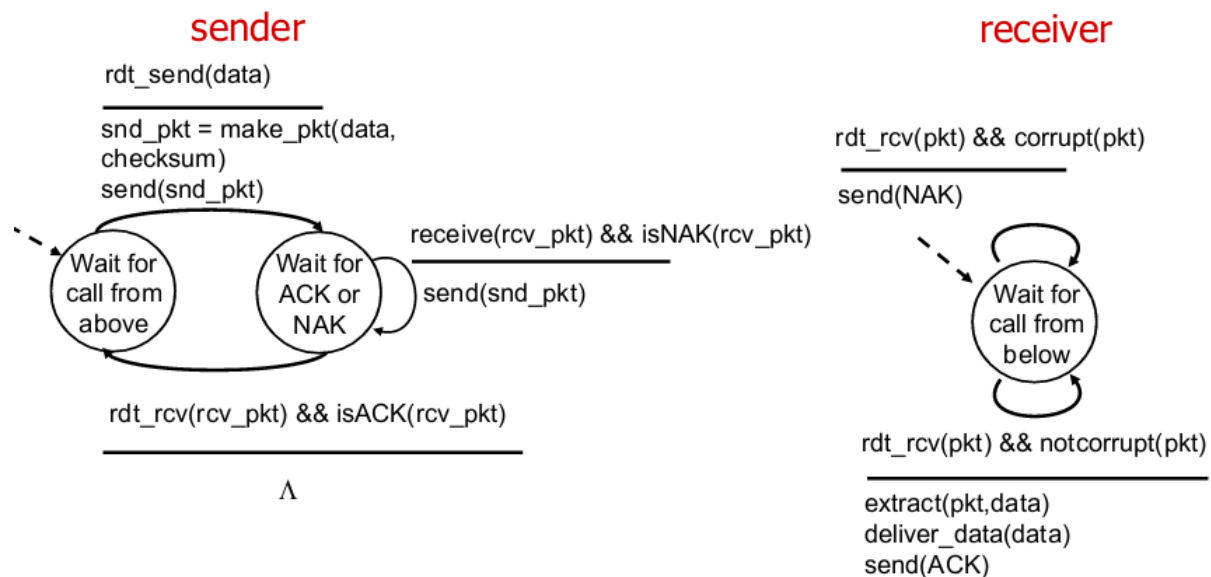
1.2 rdt2.0

1.2.1 Channel with bit errors

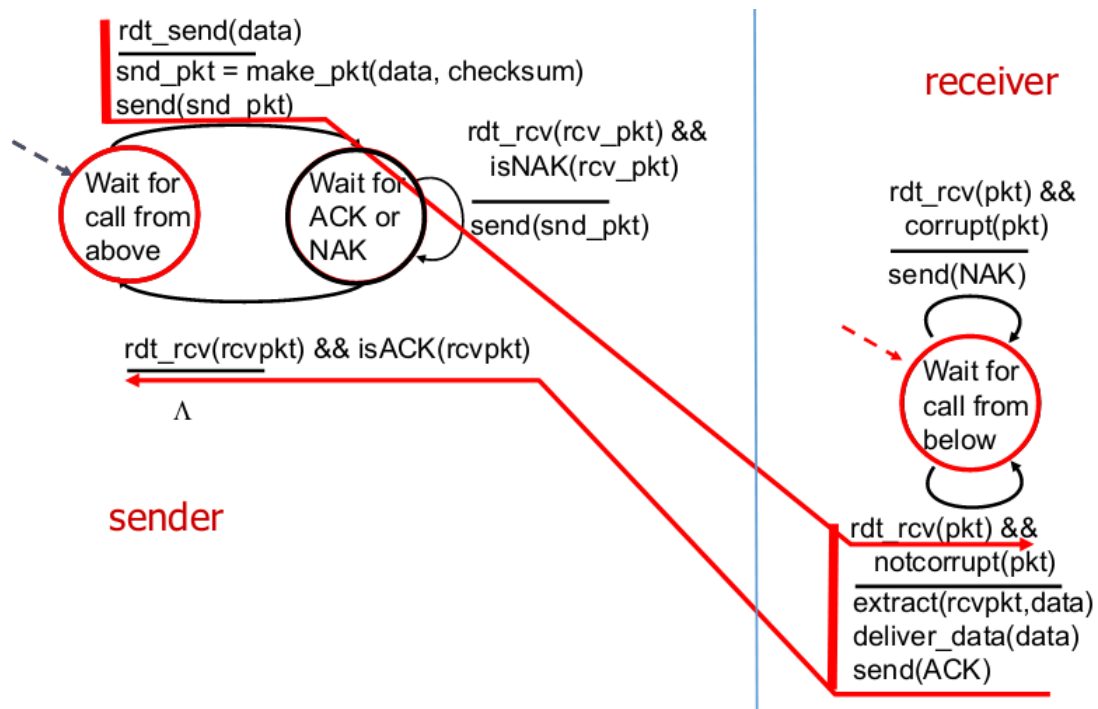
- Underlying channel may flip bits in packet, checksum to detect bit errors
- The question: how to recover from errors:
 - Acknowledgements (ACKs): receiver explicitly tells sender that pkt received OK
 - Negative acknowledgements (NAKs): receiver explicitly tells sender that pkt had errors

- Sender retransmits pkt on receipt of NAK
- Using ACKs and NAKs is known as ARQ (Automatic Repeat reQuest) protocols
 - * Error detection. Sender embeds extra bits in packets
 - * Feedback. Receiver provide sender with feedback
 - * Retransmission. Retransmit erroneous packets
- New mechanisms in rdt2.0 (beyond rdt1.0)
 - Error detection
 - Feedback: control msgs (ACK (1), NAK(0)) from receiver to sender

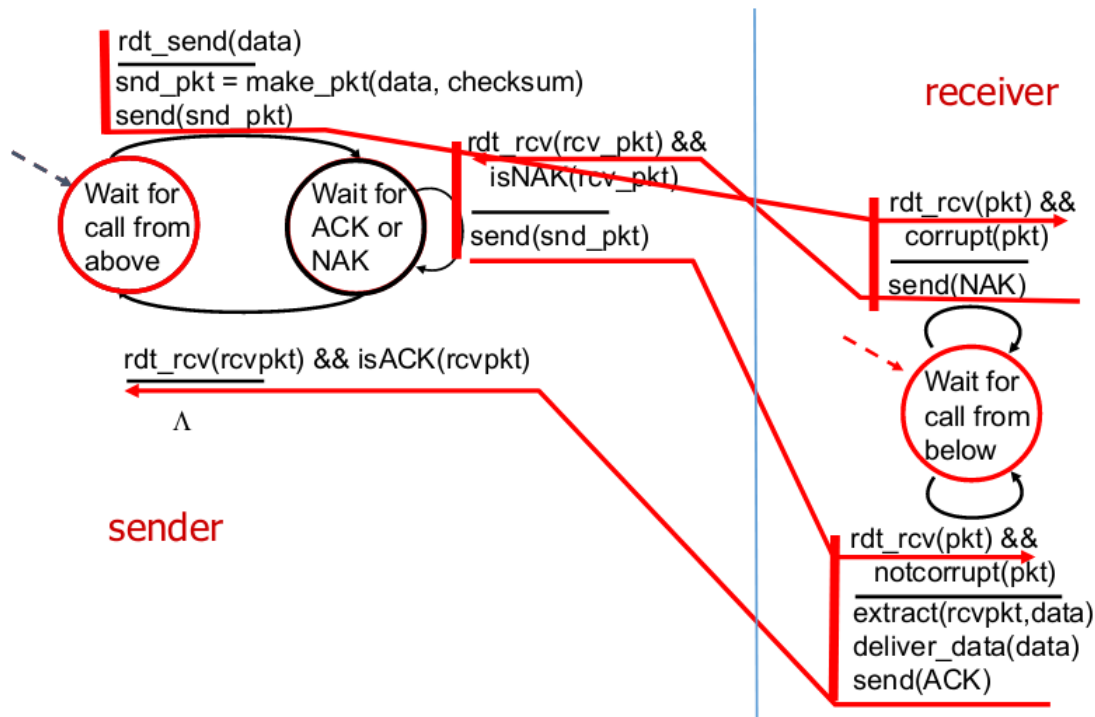
1.2.2 FSM specification



1.2.3 Operation with no errors



1.2.4 Error scenario



1.2.5 Fatal flaw

What happens if ACK/NAK corrupted:

- Sender doesn't know what happened at receiver
- Can't just retransmit: possible duplicate

Handling duplicates:

- Sender retransmits current pkt if ACK/NAK corrupted
- Sender adds sequence number to each pkt
- Receiver discards (doesn't deliver up) duplicate pkt

Stop and wait:

- Sender sends one packet, then waits for the receiver response

1.3 rdt3.0

1.3.1 Channels with errors and loss

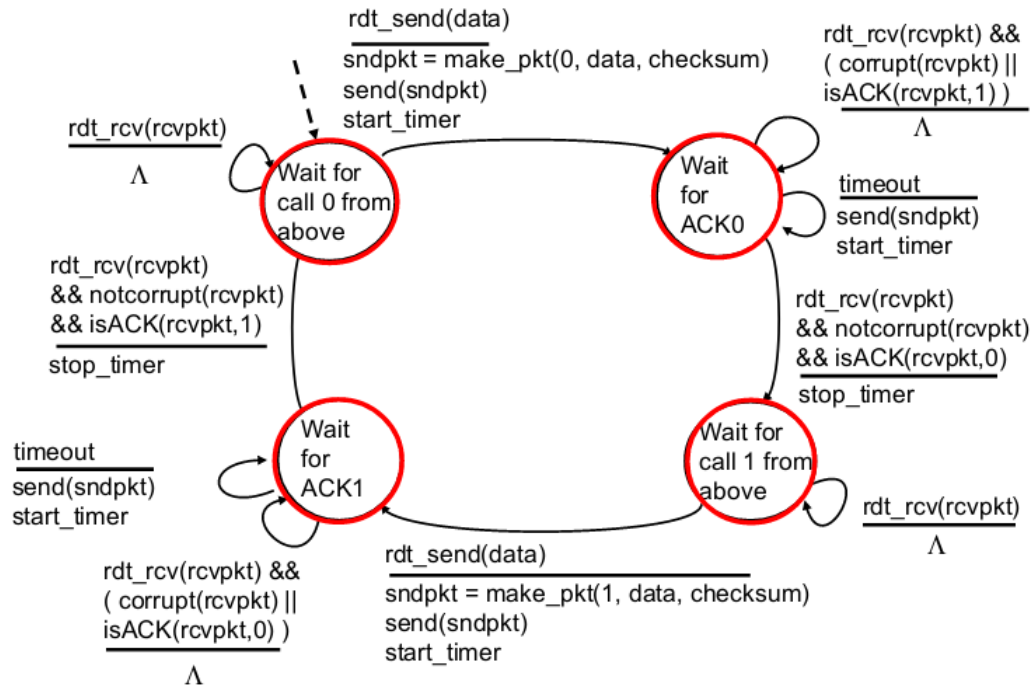
New assumption:

- Underlying channel can also lose packets (data, ACKs)
 - Checksum, seq. #, ACKs, retransmissions will be of help, but not enough

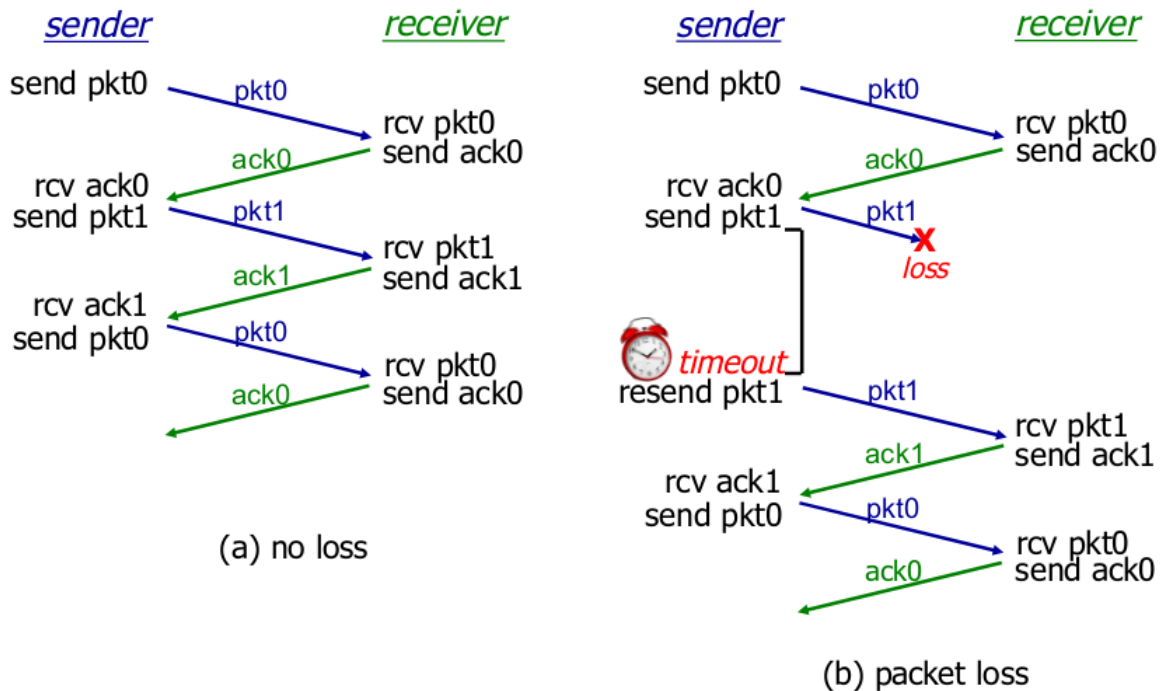
Approach:

- Sender waits "reasonable" amount of time for ACK
 - Retransmits if no ACK received in this time
 - If no pkt (or ACK) just delayed (not lost):
 - * Retransmission will be duplicate, but seq. #'s already handles this
 - * Receiver must specify seq # of pkt being ACKed
 - Requires countdown timer

1.3.2 Sender

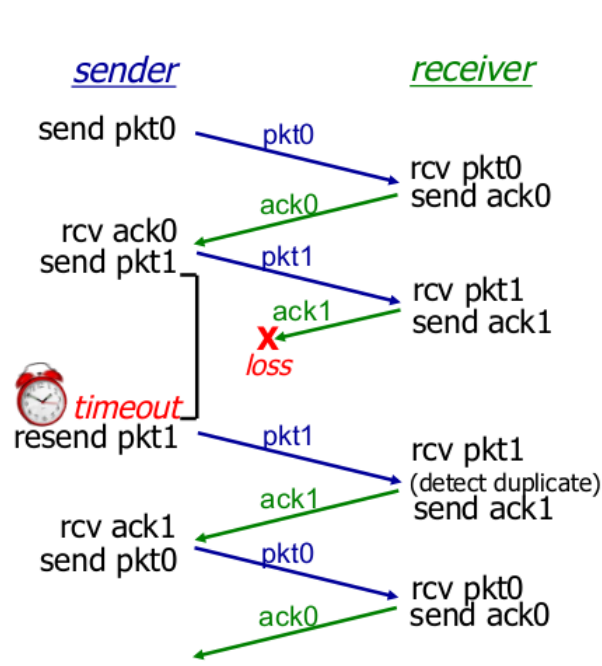


1.3.3 In action

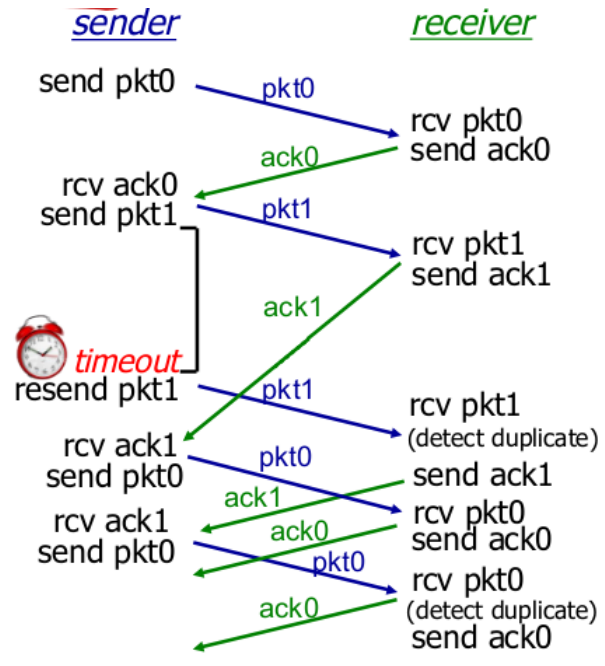


Note: alternating-bit protocol: packet sequence numbers alternate between 0 and 1.

14




(c) ACK loss



(d) premature timeout/ delayed ACK

2 Pipelined protocols



(a) a stop-and-wait protocol in operation

(b) a pipelined protocol in operation

Pipelining has the following consequences for reliable data transfer protocols:

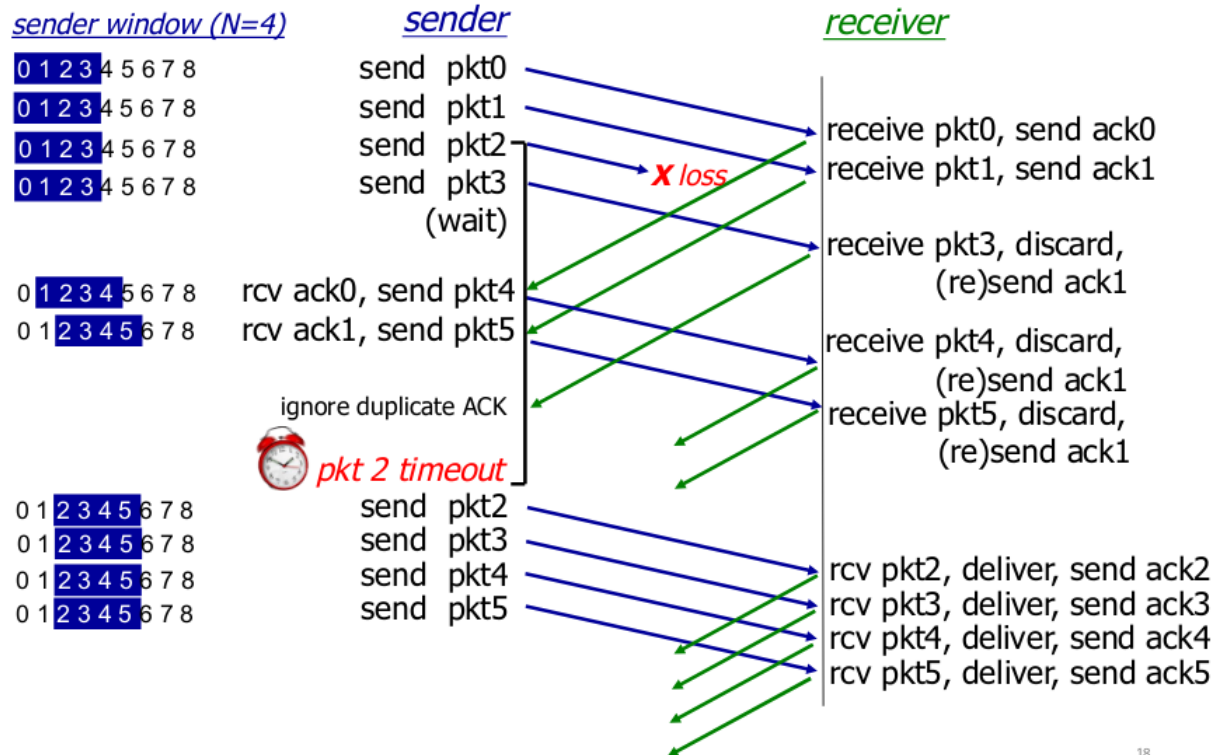
- Range of sequence numbers must be increased
 - Unique sequence number and there may be multiple, in-transit, unacknowledged packets
- Multiple packet buffering at sender and/or receiver
 - Sender buffers packets that have been transmitted but not yet acknowledged
 - Buffering of correctly received packets
- Range of sequence numbers needed and the buffering requirements will depend on the manner in which a data transfer protocol responds to lost, corrupted, and overly delayed packets

Two generic forms of pipelined protocols:

- Go-back-N
 - Sender can send multiple packets without waiting for ACK
 - Sender can have up to N unacked packets in pipeline
 - Receiver only sends cumulative ack, doesn't ack packet if there's a gap

- Sender has timer for oldest unacked packet, when timer expired, retransmit all unacked packets
- Selective repeat:
 - Sender can have up to N unacked packets in pipeline
 - Receiver sends individual ack for each packet
 - Sender maintains timer for each unacked packet, when timer expires, retransmit only that unacked packet

3 GBN in action



18

4 Selective repeat

- Receiver individually acknowledges all correctly receives pkts. Buffers pkts, as needed, for eventual in-order delivery to upper layer
- Sender only resends pkts for which ACK not received. Sender timer for each unACKed pkt
- Sender window
 - N consecutive seq #'s
 - Limits seq #'s of sent, unACKed pkts

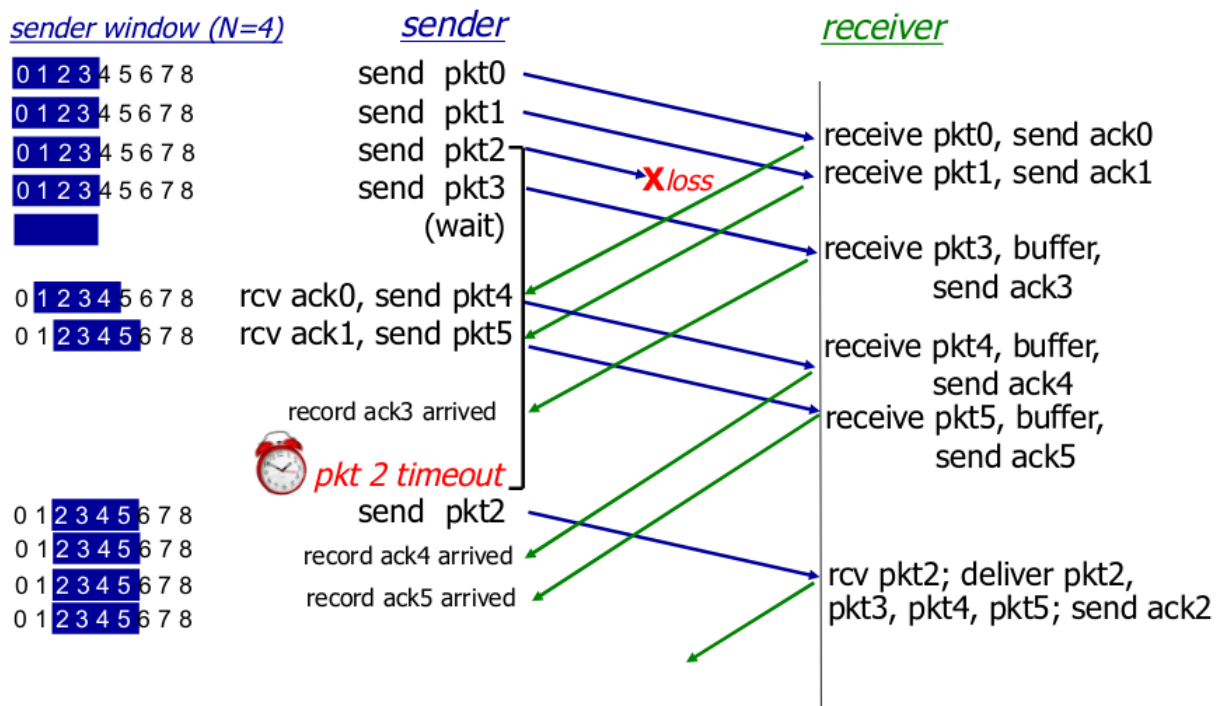
Sender:

- Data from above:
 - If next available seq # in window, send pkt
- timeout(n)
 - Resend pkt n, restart timer
- Mark pkt n as received
- If n smallest unACKed pkt, advance window base to next unACKed seq #

Receiver:

- pkt n in $[rcvbase, rcvbase+N-1]$
 - Send ACK(n)
 - Out-of-order: buffer
 - In-order: deliver(also deliver buffered, in-order pkts), advance window to next not-yet-received pkt
- pkt n in $[rcvbase-N, rcvbase-1]$
 - Ack(n)
- Otherwise:
 - Ignore


4.1 Selective repeat in action



4.2 Selective repeat dilemma

Example:

- Finite range of seq # s: 0,1,2,3
- Window size=3
 - Receiver sees no difference in two scenarios
 - Duplicate data accepted as new in (b)



receiver can't see sender side.
receiver behavior identical in both cases!
something's (very) wrong!

