

Local Search

Often in search problem the path to a goal state is of no consequence, it is the goal state itself that is the concern. For example in the 8-Queens problem it is the final configuration that is required

AI search version:

- Given a configuration of queens
 - One in each column
- Move each queen in its column
- So that a configuration of non-attacking queens is obtained

Such problems are abundant in real life

- in the scheduling of computational jobs
- in facility floor layout
- in automatic programming

If the path to a solution is of no consequence then local search algorithms

- Using only the current state and moving only to successors of that state
- Might be better employed than "global path-based" search methods

Local searches often have two advantages

- They use very little memory - having only to remember the current state
- They often give reasonable solutions in very large state space - where more systematic search algorithms are unsuitable

Definition: "Local state-based" search problem

Consists of:

- State space
 - States; state-to-state transitions; initial state
- Objective function f

Local search algorithms are useful for solving optimisation problems

- Aim is to find an optimal state according to some objective function defined on the states

A local search algorithm is optimal if:

- Whenever it finds a solution then it is a global minimum/maximum

1 Hill-climbing search

Hill climbing:

- Iteratively move to a better successor state
 - i.e., a successor state whose (objective function) f -value is higher until no such successor state exists when the algorithm terminates

Listing 1: Hill-climbing

```

1 current=initial state
2 loop do
3     successor=successor of current with highest f-value
4     if fsuccessor ≤ f(current) then return current
5     else current = successor

```

Notes:

- Hill-climbing doesn't look beyond the immediate neighbourhood
- Can get stuck in maxima/minima

2 Hill-climbing with the 8-Queens Problem

Local search algorithms include global information within the state

For the 8-Queens Problem this amounts to each state being a description of exactly where each queen is currently located

- The transition function details all states obtained by moving a single queen to another square within the same column
 - so each state has $8 \times 7 = 56$ successors
- The objective function f (to be minimized) is the number of pairs of queens that are mutually attacking each other

3 Exiting local maxima/minima - simulated annealing

Essential to this algorithm is a schedule which dictates the "rate of cooling"

- the schedule is a list which details the temperature at any given time

The algorithm iteratively performs a local search from the current state X until the temperature drops to 0 at which point the algorithm halts

A potential successor state Y in an interaction is chosen at random

- if $\Delta E = f(Y) - f(X) \geq 0$ then Y becomes the current state (f is the objective function that we are trying to maximise)

However

- If $\Delta E < 0$ then there is a chance that the successor state Y might still be chosen to become the current state (enables us to get out of local maxima)

4 Simulated Annealing

The probability that Y becomes the current state depends on

- T - the current temperature
- ΔE - the difference between the two objective values

In essence

- The "cooler" the temperature T , the less likely it is that Y is chosen
- The smaller the different ΔE , the less likely it is that Y is chosen
 - Note that ΔE is negative, so what we are saying is the worse the value $f(Y)$ in comparison to $f(X)$ the less likely that Y is chosen

These intuitive directions are encapsulated in the probability function

$$e^{\Delta E/T} = 1/e^{|\Delta E|/T}$$

A commonly used "cooling schedule"

- $\text{schedule}[1]$ is user defined and $\text{schedule}[t] = \text{schedule}[t-1]/\beta$ where β is user-defined

Listing 2: Simulated-annealing

```

1 current=initial state
2 for t=1 to ∞ do
3   T[schedule[t]]
4   if T ≤ ε then
5     return current
6   else
7     choose a successor state succ of current at random
8     ΔE=f(succ)-f(current)
9     if ΔE ≥ 0 then
10      current=succ
11    else
12      current=succ only with probability eΔE/T

```

5 Genetic algorithms

A genetic algorithm to solve a ("state-based search") problem

- starts with a randomly generated population of individuals, each with individual
 - ordinarily represented as a string of symbols
 - encoding a particular solution to the underlying problem

Each (potential) individual in the population is rated according to a fitness function

- measures how "good" an individual is (as a solution to the underlying problem)

In the 8-Queens Problem, for example

- an individual might represent a configuration of queens and be a string of digits from {1,2,...,8} of length 8
 - the first digit details the location of the first queen in the first column, the second the location of the second queen in the second column, etc.
- The fitness of an individual is the number of pairs of non-attacking queens in the configuration

Having generated an initial population P at random

- We iteratively generate a new population newP from the current one P until some appropriate terminating condition is met
 - e.g., an individual in the population is "fit enough" or the number of iterations has hit some bound

Each iteration consists of the following process repeated —P— times

- randomly select two individuals X and Y from the current population P
 - so that "fitter" individuals are more likely to be selected
- from X and Y, reproduce a child Z
- with a small probability mutate the child Z
 - by e.g. replacing one randomly-chosen symbol in the string with some randomly-chosen symbol
- Add the child Z to the new population newP
- So, after this iteration, we have that —newP— = —P—
- The current population P is now replaced with the new population newP and the next iteration begins