

# Linear Regression, Training and Loss

## 1 Linear regression

### Definition: Linear regression

A method for finding the straight line or hyperplane that best fits a set of points

$$y = b + w_1x_1$$

y - the predicted label

b - the bias, sometimes referred to as  $w_0$

$w_1$  - the weight of feature 1

$x_1$  - a feature

## 2 Training and loss

### Definition: Training a model

Learning good values for all weights and the bias from labelled examples

### Definition: Loss

The penalty for a bad prediction

### Definition: Empirical Risk Minimisation

The process of examining many examples and attempting to find a model that minimises loss

### 2.1 Squared loss

The square of the difference between the label and the prediction

$$(\text{observation} - \text{prediction}(x))^2$$

$$(y - \hat{y})^2$$

### 2.2 Mean square error

$$MSE = \frac{1}{N} \sum_{(x,y) \in D} (y - \text{prediction}(x))^2$$

(x,y) is an example where

- x is the set of features used by the model to make predictions
- y is the example's label

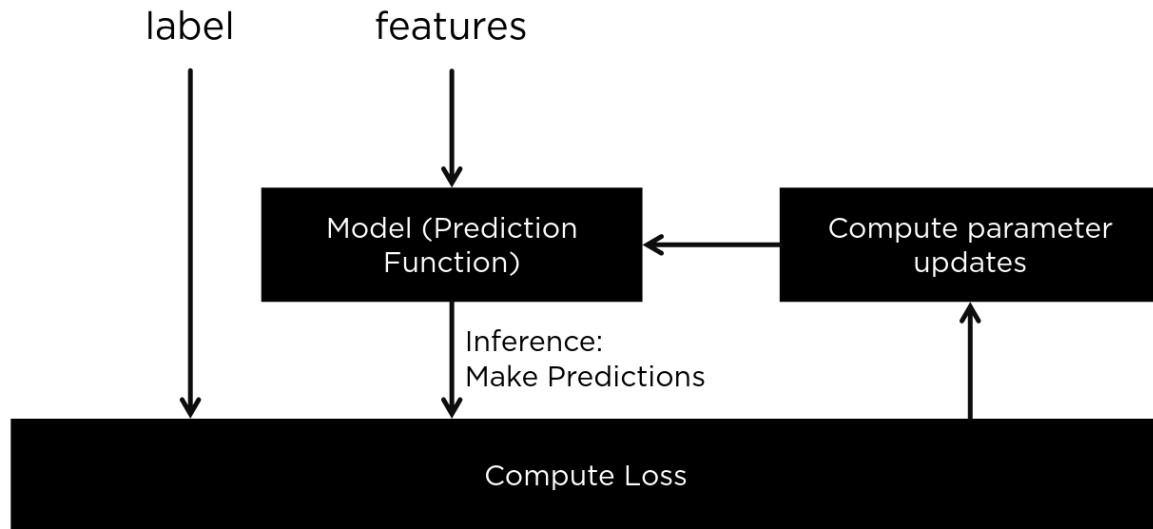
prediction(x) is a function of the weights and bias in combination with the set of features x

D is the dataset containing many labelled examples

N is the number of examples in D

### 3 Reducing loss

- Hyperparameters are the configuration settings used to tune how the model is trained
- Derivative of loss with respect to weights and biases tells us how loss changes for a given example
- So we repeatedly take small steps in the direction that minimises loss, we call these **Gradient steps**



#### 3.1 Weight initialisation

For convex problem, weights can start anywhere forming a graph that looks like  $x^2$

Foreshadowing: not true for neural networks

- More than one minimum
- Strong dependency on initial values

#### 3.2 Efficiency of reducing loss

- Could compute gradient over entire dataset on each step, but this turns out to be unnecessary
- Computing gradient on small data examples works well
- **Stochastic Gradient Descent** - one example at a time
- **Mini-batch Gradient Descent** - batches of 10-1000

#### 3.3 Learning rate

The ideal learning rate in one-dimension is

$$\frac{1}{f(x)''}$$

The ideal learning rate for 2 or more dimension is the inverse of the Hessian (matrix of second partial derivatives)